

Machine-Learning Design of Broadband Rayleigh-Wave Carpet Cloaks under Microstructure Realisability Constraints


David Aznaurov^{*a}, Davit Piliposyan^a, Sébastien Guenneau^b

^aMathematical Modelling and Machine Learning, Institute of Mechanics NAS RA

^bDepartment of Physics, Imperial College London

*Presenting author: david@mechins.sci.am

1. The dream: a cloak for surface waves
2. The catch: exact cloaks are unbuildable
3. Idea: *let a network find the material*
4. Four ML routes to a real structure
 - ▶ Neural-field material optimisation
 - ▶ Unit-cell inverse design
 - ▶ Guided-diffusion generation
 - ▶ Latent-space search
5. Results, and how much survives fabrication




A cloak made of *one* material, patterned cell by cell.

The goal

- ▶ A **Rayleigh wave** rolls along the free surface of a half-space. A surface *notch* scatters it and casts a downstream shadow.
- ▶ A **carpet cloak** is a coating that makes the surface wave behave as if the notch were not there.
- ▶ Applications: seismic protection, surface-acoustic-wave devices, non-destructive testing.

➤ **Takeaway.** We want a coating, buildable from *ordinary materials*, that restores the surface wavefield beyond the defect.



Half-space, PML boundaries, a time-harmonic point source, and the triangular cloak over the notch.

Objective

Cast carpet-cloak design as a **PDE-constrained optimisation** solved by a coordinate **neural field** through a **differentiable finite-element** wave solver, and carry the design all the way to a fabricable single-material microstructure.

- ① **Neural-field material optimisation** – the main pipeline (Sec. 3).
- ② **Unit-cell inverse design** with a neural field.
- ③ **Guided-diffusion** generation of unit cells.
- ④ **Latent-space search** over a learned design manifold.

Unconstrained design $\rightarrow \eta = 0.99$ (near-ideal)

Best *fabricable* design $\rightarrow \approx 97\%$ of ideal

First microstructured Rayleigh-wave cloak from a common material.

A coordinate map $x = \Xi(X)$, $F = \nabla_X x$, $J = \det F$, rewrites Navier's equation $\nabla_x \cdot (c^{eff} : \nabla_x u) = \rho^{eff} u_{tt}$ with the **Brun–Guenneau–Movchan** push-forward:

$$c_{ijkl}^{eff} = J^{-1} C_{IjKl} F_{iI} F_{kK}, \quad \rho^{eff} = \rho J^{-1}.$$

The obstruction (rigid statement)

The transformed tensor keeps the *major* symmetry but breaks the *minor* ones,

$$c_{ijkl}^{eff} = c_{klij}^{eff}, \quad c_{ijkl}^{eff} \neq c_{jikl}^{eff}.$$

Hence the ideal cloak is a **polar (Cosserat)** medium with a non-symmetric stress – it lies *outside* the class of ordinary Cauchy materials.

➤ **Takeaway.** The exact cloak is a theorem, not a blueprint. Our job is the *best realisable approximation* inside the Cauchy class.


Why triangular? Two phases instead of a singularity

Pinched-carpet shear map (Chatzopoulos 2023). Its gradient is *piecewise constant*:

$$F = \begin{bmatrix} 1 & 0 \\ F_{21} & F_{22} \end{bmatrix}, \quad F_{21} = \text{sign}(X_1) \frac{a}{c}, \quad F_{22} = \frac{b-a}{b} = J,$$

so the effective material is **uniform in each half** of the cloak and $\rho^{\text{eff}} = \frac{b}{b-a} \rho$ is constant.

- ▶ No inner-boundary singularity (unlike radial blow-up cloaks).
- ▶ Ideal cloak = just **two** mirror-image phases. . .
- ▶ . . . but *three* minor-symmetry pairs are broken – exactly what symmetrisation averages away by hand.



Cartesian 14×10 cell decomposition; only 46 cells fall inside the cloak (the active design variables).

Design class: orthotropic Cauchy (4CM)

Square D_2 -symmetric cells force $C_{16} = C_{26} = 0$:

$$C_{4CM} = \begin{bmatrix} C_{11} & C_{12} & 0 \\ C_{12} & C_{22} & 0 \\ 0 & 0 & C_{66} \end{bmatrix} \succ 0, \quad (+\rho).$$

Four moduli + density *per cell* are the design variables.

Objective: phase-invariant magnitude error

$$\mathcal{L}_{cloak}(\theta) = \frac{1}{|\mathcal{B}|} \int_{\mathcal{B}} \left(\frac{|u_{cloak}|}{|u_{ref}|} - 1 \right)^2 dA$$

over a near-surface strip $\mathcal{B}$ downstream of the cloak.

$$\text{Report card: the cloak ratio } \eta = \frac{\langle |u| \rangle}{\langle |u_{ref}| \rangle} \rightarrow 1 \quad (\text{perfect cloak}).$$

► **Takeaway.** Magnitudes, not complex fields: a wave routed around the notch is allowed to arrive with a benign phase lag.

① Represent


- ▶ A coordinate MLP Φ_θ is the material field.
- ▶ Outputs (C_{4CM}, ρ) per cell.
- ▶ Smoothness set by Fourier features.

② Solve & differentiate

- ▶ Frequency-domain JAX-FEM wave solve.
- ▶ Gradients via the implicit **adjoint** of the linear solve.
- ▶ End-to-end $\partial \mathcal{L} / \partial \theta$.

③ Realise

- ▶ Project onto a **realisability manifold**.
- ▶ Fill each cell with a solid/void microstructure.
- ▶ Inverse design / diffusion / latent search.



① Method 1 – Neural-field material optimisation

Reparameterise the material field by a coordinate MLP


$$\Phi_{\theta} : (x, y) \mapsto (C_{4CM}(x, y), \rho(x, y)),$$

applied as a **multiplicative residual** around a physical prior:

$$C_{4CM} = C_{4CM}^{(0)} \odot (1 + \epsilon \Phi_{\theta}^{(C)}).$$

- tanh MLP, 4 layers, Fourier positional encoding.
- Relative perturbation is scale-uniform across moduli.
- Smoothness is *built in* – no neighbour penalty needed.

Neural reparameterization of a 2D Rayleigh-wave cloak



Coordinates $\rightarrow$ Fourier encoding $\rightarrow$ MLP $\rightarrow (C_{4CM}, \rho) \rightarrow$ FEM $\rightarrow$ loss, differentiated back to θ .

► **Takeaway.** The network is the material. Design = training its weights against a wave-physics loss.

① Admissible by construction: the decoding map

Raw outputs $r = (r_{11}, r_{22}, r_{66}, r_{12})$ are decoded so every design is a valid Cauchy tensor – no penalty policing after the fact.

Positivity

$$C_{ii} = C_{ii}^{(0)} e^{\epsilon r_{ii}}$$

Log-space $\Rightarrow C_{ii} > 0$ for any real r_{ii} .

Bounded anisotropy

$$g_b = \frac{1}{2} \log R \tanh\left(\frac{g}{\frac{1}{2} \log R}\right)$$

confines $C_{11}/C_{22} \in [1/R, R]$, $R = 30$.

PD coupling

$$C_{12} = \rho_c \sqrt{C_{11} C_{22}}$$

$\rho_c = \kappa \tanh(r_{12} + \phi_0)$, $\kappa = 0.99$
 $\Rightarrow \det > 0$.

➤ **Takeaway.** A smooth, everywhere-differentiable surjection onto the set of positive-definite, bounded-anisotropy tensors: the optimiser can never leave the realisable class.


Results I – one material already beats symmetrisation

Cloak ratio η at $f^\star = 2.0$

Configuration	η
Obstacle (no cloak)	0.37
Symmetrised ideal (Chatz. 2023)	0.63
Single optimised 4CM	0.86
2×2 (four materials)	0.96
14×10 grid (46 cells)	0.99
Ideal continuous polar c^{eff}	0.999

Depth beats width: subdividing in depth helps far more than along propagation.

Surface-wave cloaking, 14×10 cells — $|\text{Re}(u)|$ ($f^\star = 2.0$)




Results II – broadband: one design, a whole octave

Train a single neural field against a band of frequencies at once:

$$\mathcal{L}_{band}(\theta) = \frac{1}{\sum_f w_f} \sum_{f \in S} w_f \mathcal{L}_{cloak}(\theta; f).$$

- ▶ Band $f^* \in [1, 3]$, 7 frequencies, equal weights.
- ▶ Optional min–max variant flattens the worst case.
- ▶ Cost scales linearly in $|S|$; frequencies solved in parallel.
- ▶ Floquet–Bloch check: propagating modes, *no* resonant band-gap trickery.



$\eta(f)$: broadband design (blue) holds ≈ 0.96 across the shaded training band; degrades gracefully outside it.

➤ **Takeaway.** The single-frequency narrowness was a limit of the *objective*, not the material class.

The gap

Optimised cells carry a homogenised (C_{4CM}, ρ) that need not correspond to any manufacturable object. We must *realise* each cell as a real two-phase (solid/void) microstructure.

- ❶ **Dataset:** $N \sim 10^6$ binary 50×50 cells (cellular-automaton growth).
- ❷ **Homogenise:** periodic FEM $\rightarrow (C_{4CM}, \rho_{eff})$; D_2 symmetry $\Rightarrow$ exactly the 4CM form.
- ❸ **Prior:** fit a GMM; keep a log-density super-level set as the realisable manifold $\mathcal{M}$.
- ❹ **Constrain:** add the manifold penalty during optimisation.

Manifold penalty

$$\mathcal{L}_{GMM} = \sum_c \max(0, \tau - \log p_{GMM}(C_{4CM}^c, \rho_c))$$

pulls every cell into $\mathcal{M}$ without pinning it to one catalogue entry.


② Method 2 – Unit-cell inverse design (neural field)

Given a per-cell target $y = (C_{4CM}, \rho)$, parameterise a *continuous density field* over the 50×50 cell, $\hat{\rho}(x, y) = g_{\phi}(x, y) \in [0, 1]$, and minimise

$$\min_{\phi} \left\| \mathcal{H}(\hat{\rho}_{\phi}) - y \right\|^2 + \lambda \mathcal{R}_{conn},$$

where $\mathcal{H}$ is the differentiable periodic-FEM homogenisation operator.

- ▶ Soft-to-hard schedule drives $\hat{\rho} \rightarrow \{0, 1\}$.
- ▶ Connectivity penalty $\mathcal{R}_{conn}$ kills floating islands/voids.
- ▶ **Deterministic**: one locally optimal cell per target.



Inverse-design realisation: one cement phase, cell-varying void pattern tiled across the carpet.

③ Method 3 – Guided-diffusion generation


Train a **conditional denoising diffusion** model on the catalogue,

$$\min_{\theta} \mathbb{E}_{t, x_0, \epsilon} \|\epsilon - \epsilon_{\theta}(x_t, t, y)\|^2, \quad y = (C_{4CM}, \rho),$$

then sample the reverse process $p_{\theta}(x_{t-1} | x_t, y)$. Because sampling is stochastic, draw N candidates and keep the closest:

$$x^{\star} = \arg \min_{n=1..N} \|\mathcal{H}(x^{(n)}) - y\| \quad (\text{best-of-}N).$$

- Prior keeps every sample *on* the manufacturable manifold.
- Explores a diverse set of realisations per target.



Best-of- N diffusion realisation of the same target field – visibly different, equally manufacturable patterns.

④ Method 4 – Search in a learned latent space


Direct pixel optimisation is ill-conditioned ($\sim 10^4$ DoFs, bad minima). Train an **autoencoder** on $(C, \rho, f_\star, \mathcal{L})$ with a *frequency-separated* latent

$$z = z_{sh} + z_f(f_\star),$$

(geometry vs. per-frequency offset), a performance head, and a within-frequency contrastive ranking on z_{sh} . Then **freeze** the AE and optimise the *code* z :

$$\min_z \alpha \max_i \mathcal{L}(f_i; z) + (1 - \alpha) \frac{1}{N} \sum_i \mathcal{L}(f_i; z), \quad \alpha : 1 \rightarrow 0.$$

► **Takeaway.** The latent search inherits the training distribution's manufacturable priors – smoother layouts, fewer iterations than raw topology optimisation.




AE latent decodes to structured, *manufacturable* motifs. Cf. `frequency_sweep_vis.png`: AE $\eta \approx 0.965$, neural-field ≈ 0.985 .

Results III – how much survives fabrication?

Realising the 20×15 cloak (100 cells), $f^\star = 2.0$

Realisation of the 100 cells	η
Continuous 4CM optimum	0.99
Nearest-neighbour snap	0.82
Inverse design (neural field)	0.94
Diffusion (best-of-N)	0.97
Ideal continuous polar c^{eff}	0.999



$\eta(f)$ across the band: diffusion (green) stays closest to the perfect-cloak line; smallest mean distortion.

► **Takeaway.** Better per-cell matching $\Rightarrow$ better cloak: best-of- N diffusion recovers **97%** of the ideal, from fabricable single-material cells.

- ▶ A single **differentiable pipeline** – coordinate neural field + JAX-FEM adjoint – designs Rayleigh-wave carpet cloaks inside the realisable Cauchy class.
- ▶ **Near-perfect** single-frequency ($\eta = 0.99$) and **broadband** (≈ 0.96 over an octave) cloaking, beating closed-form symmetrisation.
- ▶ Four ML routes to a real structure – neural-field optimisation, inverse design, guided diffusion, latent search – with **diffusion recovering** $\approx 97\%$ of ideal performance from a *single* constituent.

Future work: a differentiable generative manifold queried *during* optimisation; multiscale generation with connected neighbours on a body-fitted mesh; 3-D cloaks; and polar/Cosserat microstructures to close the broken-minor-symmetry gap.

- [1] Brun, Guenneau, Movchan (2009). *Achieving control of in-plane elastic waves*. Appl. Phys. Lett.
 - [2] Norris, Shuvalov (2011). *Elastic cloaking theory*. Wave Motion.
 - [3] Nassar, Chen, Huang (2018). *Polar metamaterials for elastic cloaking*. (Cosserat lattice).
 - [4] Chatzopoulos et al. (2023). *Symmetrised Rayleigh-wave carpet cloaks*. (baseline).
 - [5] Wang et al. (2022). *Database-selected orthotropic mechanical cloaks*.
 - [6] Tancik et al. (2020). *Fourier features let networks learn high-frequency functions*.
 - [7] Nakarmi et al. (2024). *Cellular-automaton generation of mesostructure libraries*.
 - [8] Yang et al. (2024). *Guided diffusion for inverse microstructure design*.
 - [9] Xu et al. (2020). *JAX-FEM: a differentiable finite-element solver*.
- (Placeholder entries – replace with the paper's `cas-refs.bib` for the final deck.)

Thank You

Acknowledgements

The authors gratefully acknowledge Danila Rukhovich (Google) for support. Code, dataset, and the differentiable JAX-FEM pipeline available on GitHub <https://github.com/M3LLab/nn-cloaking>

